

## Future Thoughts: On Facetwork's Positioning in the AI-Agent Era — A Dissenting Companion

1. Where the dominant framing is right
2. Where the dominant framing overshoots
  - 2.1 "Target existing engines" undersells FFL's semantics
  - 2.2 The human-DSL versus AI-IR dichotomy is too clean
3. Restatement
4. What this implies for the roadmap
5. Working evidence: the structured-diagnostics feedback loop
6. Working evidence: vendor surfaces as composable, FFL-native units

# Future Thoughts: On Facetwork's Positioning in the AI-Agent Era — A Dissenting Companion

---

A short companion to `future-thoughts-ai-native.md` and `ai-authorship.md`.

This note records one reader's partial agreement and partial pushback against the now-common framing that a new workflow system — including Facetwork — has lost most of its standalone value because AI agents can generate workflows for mature engines (Temporal, Step Functions, Prefect, LangGraph, Dapr Workflows). The dominant reframing says: *don't compete on syntax, compete on validation, policy, repair, provenance, and safe execution of AI-generated plans — possibly by targeting existing runtimes*. This note agrees with most of that, and dissents on two points.

---

## 1. Where the dominant framing is right

---

The core claim holds up. If Facetwork's pitch is "a nicer DSL so humans can author workflows more easily," the pitch is weaker in 2026 than it was in 2018. Three reasons:

- **AI agents don't reward elegant syntax.** An agent generating a Temporal workflow, a Step Functions state machine, or an Airflow DAG does not particularly benefit from a smaller grammar. Given documentation and examples, agents produce correct definitions in verbose, imperative target languages just as readily as in a compact DSL.
- **Mature engines have runtime moats that are expensive to replicate.** Durability, retries, scaling, observability, connectors, and — crucially — the operational trust accumulated over years of production usage. A new runtime cannot shortcut that.
- **The binding constraint has moved.** The hard problem is no longer "how do I express this workflow?" It is "how do I verify, bound, explain, and repair a workflow that was generated by a system I did not fully supervise?" That is a validation and control problem, not a syntax problem.

So the reframing is sound. The load-bearing pieces of Facetwork, in this light, are not the surface syntax of FFL but the machinery around it: the typed schemas, the compiler-level validation, the step recovery actions, the repair tooling, the execution traces. Those are AI-native control-plane features, whether or not they were originally framed that way.

## 2. Where the dominant framing overshoots

---

Two pushbacks.

## 2.1 "Target existing engines" undersells FFL's semantics

The natural next move from the reframing is: *let Facetwork be a frontend — an AI-native IR and validator — that compiles down to Temporal or Step Functions*. This is tempting because it inherits the runtime moat. But it flattens the thing that may actually matter.

FFL's composition model is not just syntactic sugar over Temporal's imperative workflow functions or Step Functions' state machines. Mixins, implicit facets, typed schemas, `andThen` `foreach`, statement-level `andThen`, `catch when`, prompt blocks, script blocks — these form a different *algebra*. Workflows compose by declaration, not by control flow. Effects and data dependencies are explicit in the structure, not buried in procedural code.

Why does that matter for AI-generated workflows? Because **compositional, declarative structure is easier to verify than imperative structure**. A validator can reason about a facet graph — its dependencies, its typed inputs and outputs, its declared effects — far more tractably than it can reason about an arbitrary Python function that happens to call `workflow.execute_activity()`. A smaller, more regular surface area shrinks the set of things a validator, a repair tool, or a human reviewer has to consider.

Compiling FFL to Temporal would erase that advantage. The target runtime's expressiveness leaks back into the IR, and the IR becomes "whatever Temporal can do," which is "arbitrary code." The control-layer framing only works if the IR is *more constrained* than the runtime it ultimately executes on. Facetwork's semantic restrictions are a feature, not an accident to be hidden behind a compiler.

So: yes to AI-native IR. No to reducing FFL to a pretty frontend for somebody else's runtime. The semantics are part of the product.

## 2.2 The human-DSL versus AI-IR dichotomy is too clean

The reframing tends to sort workflow systems into two camps: human-facing DSLs (declining in value) and machine-facing IRs (rising in value). Facetwork, in this sorting, is pushed to pick a side and told to pick the IR side.

But the high-value artifact in an AI-agent workflow may be precisely the thing that is *both*. When an agent proposes a workflow, a human reviewer needs to read it to approve it. A pure IR — JSON, protobuf, a Merkle DAG of hashes — loses the review affordance. Reviewers will not read thousands of lines of serialized graph nodes and catch subtle misbehavior. They need a projection that is dense, readable, and structurally honest about what will execute.

A well-designed human-readable DSL, if its semantics are tight, *is* that projection. FFL is in that middle zone: readable enough that a reviewer can see the shape of what was generated, structured enough that a compiler can validate it. A companion note (`future-thoughts-ai-native.md` §1) captures this as "the language is really two things: a canonical IR (what executes) and a projection layer (what a human sees when verifying)." That is exactly the point. FFL's surface syntax is not a legacy concession to human authors; it is the review interface for AI-authored plans.

The risk of the pure-IR framing is that it dismisses the projection as nostalgia. It is not. It is the seam where human oversight actually happens.

### 3. Restatement

---

Synthesizing agreement and dissent:

- **Agreed:** Facetwork's strongest modern positioning is as an AI-native workflow specification and control layer, centered on validation, policy, repair, and provenance.
- **Dissented:** That positioning does not imply abandoning FFL's semantics in favor of existing runtimes, nor collapsing the human-readable surface into a pure machine IR. FFL's compositional algebra is part of what makes validation tractable, and its readable surface is part of what makes human oversight possible.

The sharper version of the pitch, then, is not:

Facetwork is an AI-native workflow specification and control layer that helps agents generate, validate, explain, repair, and safely execute workflows.

But rather:

Facetwork is an AI-native workflow system whose semantic restrictions make AI-generated plans tractable to validate and repair, and whose readable surface makes them tractable to review.

The claim rests on two structural properties — restricted semantics and a readable projection — that reinforce each other. Drop either one and the system loses what distinguishes it from both the mature runtimes on one side and the pure-IR proposals on the other.

### 4. What this implies for the roadmap

---

Concretely, if this framing is right:

- **Keep the semantic restrictions.** Resist feature requests that would make FFL arbitrarily expressive. Every restriction is a validator's friend.
- **Invest in the projection layer as a first-class review interface.** Execution traces, diff-against-prior-run views, natural-language paraphrases of generated plans, and effect summaries are not UI polish — they are the oversight mechanism.
- **Treat the IR and the surface as a pair.** If a proposed feature is readable but not validatable, or validatable but not readable, it probably does not belong.
- **Be skeptical of "just compile to Temporal."** That path is available, but taking it would spend the semantic advantage to buy a runtime moat Facetwork does not necessarily need at its stage.

None of this contradicts the dominant reframing. It sharpens it: the AI-native control layer works because of FFL's semantics and because of FFL's surface, not in spite of them.

## 5. Working evidence: the structured-diagnostics feedback loop

---

The positioning above is testable. If FFL's restricted semantics make validation tractable for AI authors, then we should be able to give an agent the same diagnostic-driven repair loop it already enjoys for Python or TypeScript: error → linked docs → wrong/right examples → fix → re-check. The recent validator and MCP changes are a deliberate test of that claim.

Concretely, the validator now emits structured `ValidationError` records carrying not just a message and a source location but a stable `rule_id`, a `severity`, a `docs_uri`, and an optional `suggested_fix`. Forty-six call sites across the validator have been threaded with explicit `rule_ids` spanning nine categories — reference resolution, `when` exhaustiveness, type checking, yield targets, prompt and script blocks, schema and implicit-facet placement, and structural restrictions on where workflows and schemas may appear. Every emitted `rule_id` has a paired markdown file at `docs/reference/rules/{rule_id}.md` containing a short rationale, a wrong example, a right example, and the underlying constraint it enforces. The set is currently exact: 41 emitted `rule_ids`, 41 documented `rule_ids`, with a script that diffs the two on every change.

The MCP layer exposes this without flattening it. A call to `fw_validate` returns the full diagnostic structure — `rule_id`, `severity`, `location`, `docs_uri`, `suggested_fix`, plus warnings — rather than a single English error string. The companion resource `afl://docs/rules/{rule_id}` returns the rule's documentation, so an agent that hits a diagnostic can fetch its own remediation context in one step. A second resource exposes a small set of canonical, validator-clean FFL examples, so the agent has a known-good starting point when the failing source is too tangled to repair in place.

Two things are worth noting about why this works at all, and they come straight from the dissent above.

First, the rule taxonomy is small and stable because the language is small and stable. The validator emits 41 distinct rules, not 410. That number is finite for the same reason FFL is restricted: the grammar does not let workflows or schemas float free of a namespace; references must take `step.field` form; `when` blocks must end in a default; loop variables are scoped; comparisons return Boolean and `&&` / `||` / `!` accept only Boolean operands. Each of those restrictions is an axis along which the validator can speak precisely. A more permissive language would not have a closed enumeration of failure modes — it would have a long tail of "here's a parser error, good luck." The compositional algebra in §2.1 is what makes the `rule_id` space tractable to enumerate and tractable to document.

Second, the suggested fixes and the docs uri are valuable to a *human reviewer* of agent-generated FFL, not only to the agent. A reviewer reading a generated workflow can click through to `afl://docs/rules/REF_INVALID_STEP_FORMAT` and see the same wrong/right framing the agent saw. The diagnostic surface is, in this sense, another instance of the projection layer argued for in §2.2: the same artifact serves machine repair and human verification. A pure-IR system would have to invent a separate review affordance for this; FFL inherits it because the surface and the semantics are paired.

This does not, on its own, prove the larger positioning. What it does demonstrate is that one of the load-bearing claims — *restricted semantics make AI-driven repair tractable in the same way restricted semantics make human review tractable* — is not aspirational. It is now the working contract between the compiler and any agent (Claude Code, an MCP client, a future planning agent) that touches FFL. Two further claims flagged here have since been built on the same template — a small, restricted vocabulary surfaced to both agent and reviewer. **Effect annotations on event facets** now exist (with `Effect(kind=...)` / `with Cost(tier=...)` mixins, surfaced and filterable through the `fw_capabilities` capability index; all 247 facets of `fw_h_osm` annotated): a policy check, or an authoring agent choosing between a cheap pure primitive and an expensive external one, finally has something concrete to read — the restricted-vocabulary-plus-discovery shape, applied to effects. The remaining runtime-layer claim is structured repair diagnostics for stuck workflows (the analogue at the execution layer of what the validator does at the source layer); it follows the same template and extends the same evidence.

## 6. Working evidence: vendor surfaces as composable, FFL-native units

---

The dissent's §2.1 claim is that FFL's compositional algebra would be erased if Facetwork were merely a frontend that compiled down to Temporal or Step Functions. The recent `fw_h_anthropic` package is a working test of that claim from the opposite end: rather than compiling out, it composes *in* — taking an external vendor surface (the public APIs at [github.com/anthropics](https://github.com/anthropics)) and exposing it as 16 typed event facets across six FFL namespaces (Messages, Batch, Files, Agent SDK, Claude Code, Computer Use), distributed as a standalone pip-installable package discoverable through the `facetwork.examples` entry point.

The package is small — six area directories, one CLI per facet, one short `_lib/<area>.py` per area — but the structural point is large. The unit of vendor-wrapping is FFL, not Python. `anthropic.messages.CreateMessage` has a typed signature (`prompt: String, system: String = "", model: String = "", max_tokens: Long = 1024, cache_system: Boolean = false`) and a typed return (`MessageResult`). Any FFL workflow can call it by name; any handler implementation can be substituted at registration time; any reviewer can read it without reading Python. The same is true for `anthropic.batch.RunBatch`, for `anthropic.files.UploadFile`, for the cross-area `anthropic.compose.DocumentQA` workflow that bundles upload + question into one typed step. None of these would render naturally as activities in a Temporal `@workflow` class or as states in a Step Functions JSON document, because in both of those models a vendor wrapper is *just code*, indistinguishable from any other Python or JSON.

In Facetwork it is something else: a *facet catalog*. A reviewer can read it as a contract. A workflow author can compose against it without importing it. An operator can swap its implementation (real API, mocked client, recorded fixture) without touching the FFL that uses it. A second package can publish a workflow that consumes facets from the first by qualified name, exactly as if both lived in the same monorepo. The `research-agent` example was ported to this shape: it consumes `anthropic.messages.CreateMessage` from `fwh_anthropic` and parses the response through small typed `Parse<Schema>` event facets local to its own namespace. From the consumer's FFL, there is no syntactic marker that the LLM call lives in a separate package; from the operator's perspective, that separation is what lets the two packages move on independent release cadences without breaking each other.

This is the dissent's §2.1 made concrete. The compositional algebra — qualified naming, typed signatures, handler registration through a database, entry-point discovery, cross-namespace `use imports` — is the substrate that makes vendor wrapping cleanly cross-package. A frontend that compiled down to Temporal would not have this property; the moment FFL became "valid Temporal activity definitions" the qualified-naming game would be over and `fwh_anthropic` would have to ship as a Python library that exports activity functions rather than as a package that *publishes facets*. The vendor-wrapping use case is one of the more visible places where the difference between *facet catalog* and *callable library* shows up in the operational shape.

Two secondary observations belong to the same evidence cluster.

First, the JSON-bridge fields that `fwh_anthropic` uses for nested-list payloads (`tools_json`, `messages_json`, `trace_json`, `results_json`, `files_json`) are an honest acknowledgement that today's FFL type system does not model arbitrary lists of records as values. The fields are not elegant. They work because the surrounding compositional algebra still does — the bridge crosses one layer of opacity (the JSON-shaped payload) but preserves typed parameters and typed returns at the facet boundary. A future FFL extension that added typed records-in-lists would eliminate these bridges without changing any of the cross-package composition above. The dissent's §2.1 argument predicts exactly this: the algebra is robust to additions that tighten the type system; it would be brittle to amputations that flatten the algebra into someone else's runtime.

Second, the runtime issues that surfaced during the `fwh_anthropic` work — three of them, detailed in thesis §14.4 — are the kind of bug a typed-DSL approach catches because the surface forces the question. `resume_step` not re-queueing stuck siblings of a dirtied block, and `get_completed_step_by_name` rejecting steps with populated returns but not-yet-complete status, are bugs that exist only because the runtime distinguishes between *a step has returns to read* and *a step has finished its lifecycle*. A purely procedural orchestrator does not have that distinction to get wrong, but it also does not have it to use: the inline-andThen pattern `g0 = Gather(...)` andThen `{ f0 = Extract(sources = g0.sources) }` is not expressible at all in a system whose composition primitives are imperative function calls. The dissent's broader point — that the load-bearing properties live in the algebra, not in the surface — is illustrated by both directions. The algebra makes hard-to-catch bugs visible, and it makes hard-to-write compositions trivial.

The earlier claim — that FFL's semantic restrictions and its readable surface are paired — extends from the source layer (the validator + docs-by-rule\_id of §5) to the package layer (vendor wrappers as facet catalogs) and to the runtime layer (composition bugs discovered through real workflow execution). Three layers, one substrate. None of this is incompatible with the "AI-native control layer" framing of §3; it is a refinement of what that framing *is*. The control layer happens to be a language, the language happens to compose vendor surfaces cleanly, and the vendor surfaces happen to be the most natural carriers for the effect typing that the broader AI-agent agenda eventually wants. The dominant reframing was right that Facetwork's value sits in the control layer. The working evidence keeps suggesting that the surface is part of the control layer, not separate from it.